Towards Operationalising Cloud Intelligence Services using Software Engineering Practices

Alex Cummaudo
BSc *Swinburne*, BIT(Hons)

Applied Artificial Intelligence Institute
Deakin University
Melbourne, Australia

October 16, 2018

Contents

List of Figures

List of Tables

List of Listings

Chapter 1

Introduction

Within the last half-decade, we have seen an explosion of cloud-based services typically marketed under an AI banner. Vendors are rapidly pushing out AI-based solutions, technologies and products that encapsulate half a century worth of machine-learning research: a 2016 report by market research company Forrester captured such growth into four key areas [72] as replicated in Figure 1.1. Moreover, developers eager to develop a next generation of software are shifting away from mobile-first to 'AI-first' apps, that will reason, sense, think, act, listen, speak and execute our whims right within the palms of our hands. Most prominently spearheading this wave of AI-first thinking is Google, as evident through their 2018 rebranding of *Google Research* to *Google AI* [52].

These services aim to lower the entry barrier to develop, test and deploy AI-first software in both skill and time. Software engineers needn't require a formal training in machine-learning nor a strong understanding of mathematics: thus, *skill required* is reduced. The training of such classifiers involves the laborious process of sourcing, curating and labelling large datasets: using such services does not, and thus *time* is reduced. To this end, they needn't require much machine-learning expertise or experience at all; instead, the process is abstracted behind an API call, only requiring knowledge on how to use a RESTful architecture [41] to access the cloud service.

To contrast this with more traditional means, a developer may choose to write up a deep-learning NN (for example) and train it using their own dataset. While this is laborious in time and demands significant knowledge in machine learning, the developer has full control over the models she creates. Alternatively, she may choose to download a pre-trained model and ML framework, such as Tensorflow [16]; less demanding in time but still requiring the knowledge to wire-up models with frameworks.

Figure 1.1: A Broad Range of AI-Based Products And Services Is Already Visible. (From [72].)

Category	Sample vendors and products	Typical use cases
Embedded AI Expert assistants leverage AI technology embedded in platforms and solutions.	• Amazon: Alexa • Apple: Siri • Facebook: Messenger • Google: Google Assistant (and more) • Microsoft: Cortana • Salesforce: MetaMind (acquisition)	• Personal assistants for search, simple inquiry, and growing as expert assistance (composed problems, not just search) • Available on mobile platforms, devices, the internet of things • Voice, image recognition, various levels of NLP sophistication • Bots, agents
AI point solutions Point solutions provide specialized capabilities for NLP, vision, speech, and reasoning.	• 24[7]: 24[7] • Admantx: Admantx • Affectiva: Affdex • Assist: AssistDigital • Automated Insights: Wordsmith • Beyond Verbal: Beyond Verbal • Expert System: Cogito • HPE: Haven OnDemand • IBM: Watson Analytics, Explorer, Advisor • Narrative Science: Quill • Nuance: Dragon • Salesforce: MetaMind (acquisition) • Wise.io: Wise Support	• Semantic text, facial/visual recognition, voice intonation, intelligent narratives • Various levels of NLP from brief text messaging, chat/conversational messaging, full complex text understanding • Machine learning, predictive analytics, text analytics/mining • Knowledge management and search • Expert advisors, reasoning tools • Customer service, support • APIs
AI platforms Platforms that offer various AI tech, including (deep) machine learning, as tools, APIs, or services to build solutions.	• CognitiveScale: Engage, Amplify • Digital Reasoning: Synthesys • Google: Google Cloud Machine Learning • IBM: Watson Developers, Watson Knowledge Studio • Intel: Saffron Natural Intelligence • IPsoft: Amelia, Apollo, IP Center • Microsoft: Cortana Intelligence Suite • Nuance: 360 platform • Salesforce: Einstein • Wipro: Holmes	• APIs, cloud services, on-premises for developers to build AI solutions • Insights/advice building • Rule-based reasoning • Vertical domain advisors (e.g., fraud detection in banking, financial advisors, healthcare) • Cognitive services and bots
Deep learning Platforms, advanced projects, and algorithms for deep learning.	• Amazon: FireFly • Google: TensorFlow/DeepMind • LoopAI Labs: LoopAI • Numenta: Grok • Vicarious: Vicarious	• Deep learning neural networks for categorization, clustering, search, image recognition, NLP, and more • Location pattern recognition • Brain neocortex simulation

Figure 1.2: Number of posts categorised on StackOverflow under popular computer vision cloud intelligence services.

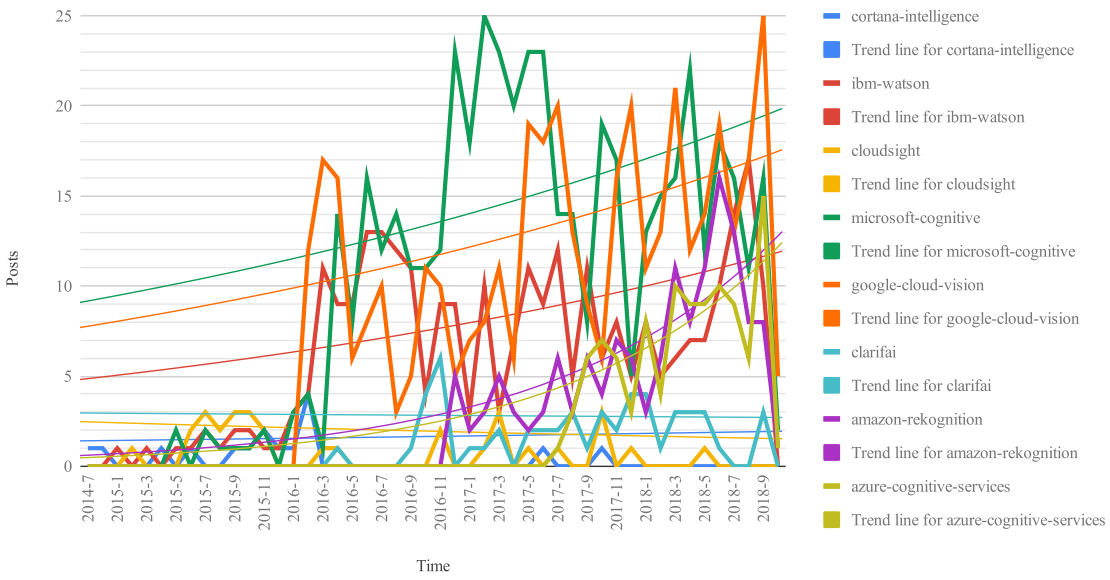

With less time and skill required to build AI-first apps using these cloud services, these services have begun to gain traction within developer circles: Figure 1.2 shows the increasing trend of posts since 2014 on StackOverflow that categorise popular computer vision cloud APIs.[1] A growing popularity into such services sparked varied nomenclature: Cognitive or Intelligence Services, Artificial Intelligence or Machine Learning as a Service [88], Cloud Machine Learning and so on. We henceforth refer to such services under the term 'Cloud Intelligence Services' (CISs), and diagrammatically express their usage within Figure 1.3.

Figure 1.3: Overview of Cloud Intelligence Services.

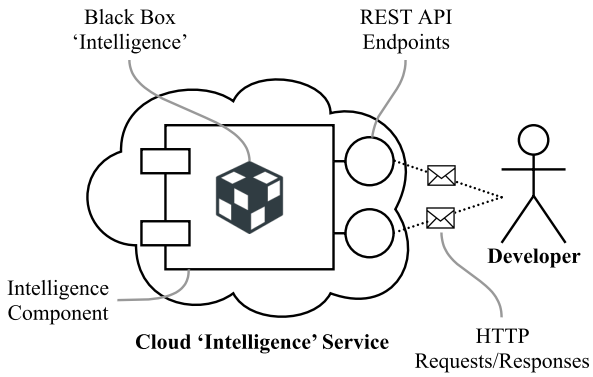

A developer accesses a CIS component via a RESTful HTTP API endpoint(s). For their given input, they receive an intelligent-like response typically formatted in JSON. We note the

[1]Query run on 12 October 2018 using StackExchange Data Explorer. Refer to `https://data.stackexchange.com/stackoverflow/query/910188` for full query.

intelligence component masks its 'intelligence' through a black-box: in recent years, there is a rise in providing human-level intelligence via crowdsourcing Internet marketplaces such as Amazon Mechanical Turk [12] or ScaleAPI [14]. Thus, a CIS may be powered by varying degrees of intelligence: human intelligence, computer manipulation or brute-force methods, or supervised or unsupervised learning.

While there are many types of CISs evident (such as OCR Transcription, Object Categorisation, Object Comparison, NLP etc.), we scope the work investigated to computer vision CIS analysers [8, 3, 1, 13, 9, 5, 4, 7, 10, 15, 11, 6, 2]. The ubiquity of computer vision CISs is exemplified through evermore growing applications that use these APIs: aiding the vision-impaired [86, 31], accounting [73], data analytics [55], and student education [34].

1.1 Motivation: Current Developer Mindsets'

Figure 1.2 shows an increasing trend to the adoption and discussion of CISs with developers. As aforementioned, these services are accessible through APIs and consist of an 'intelligence' black box (Figure 1.3). When a term 'black box' is used, the input (or stimulus) is transformed to its to outputs (or response) without any understanding of the internal architecture by which this transformation occurs; indeed, this well-understood theory derives from electronic sciences since the 1950s–60s **[cite:Bunge, Mario 1963; Ashby, W. Ross, 1956, chapter 6]**. In many cases, these black boxes are inherently probabilistic and stochastic: for instance, a computer vision CIS usually returns the *probability* that a particular object (the response) exists in the raw pixels (the stimulus), and thus we must stochastically retrieve hundreds of these results to get an interpretation of the *distribution of overall confidence* returned from the service. There are thus therefore three factors to consider when implementing and using a CIS: (i) API usability, (ii) the nature of stochastic and probabilistic systems, and (iii) how both impact on software quality.

1.1.1 The Impact on Software Quality

APIs reflect a set of design choices made by their providers. Evaluations of API usability advocate for the accuracy, consistency and completeness of APIs and their documentation [83, 91] written by providers, while providers should consider mismatches between the developer's conceptual knowledge of the API its implementation [65]. It is therefore imperative that CIS providers consider the impact of their API usability. Poor API usability, therefore, hinders

on the internal quality of development practices, slowing developers down to produce the software they need to create.

Moreover, developers need to be wary of the probabilistic nature of probabilistic systems. These APIs become inherently non-deterministic in nature, but developers are still taught with the deterministic mindset that all API calls are the same. Simple arithmetic representations (e.g., $2+2=4$) will *always* result in 4; but a multi-layer perceptron neural network performing similar arithmetic representation [22] gives the probability where the target output (*exactly* 4) and the output inferred (*possibly* 4) matches as a percentage (or as an error where it does not match). That is, instead of an exact output, there is instead a *probabilistic* result: $2+2$ *may* equal 4 with a confidence of n. External quality must therefore be considered in the outcome of these systems, such as in the case of thresholding values, to consider whether or not the inference has a high enough confidence to justify its result to end-users.

In order to fully understand this problem, there are multiple dimensions one must consider: the impact of software quality; the fact that these systems underneath are probabilistic and are stochastic; the cognitive biases of determinism in developers; the issue of consistency in API usage. While existing literature does extensively explore software quality and API usability, these studies have only had emphasis on deterministic systems and thus little work to date has investigated such factors on probabilistic systems that make up the core of computer vision CISs. We explore more of these facets in the motivating scenario below.

1.1.2 Motivating Scenario

How do developers work with a CIS? How usable are these APIs, and how well do developers understand the non-deterministic and stochastic nature of a deep-learning cloud-based API? To motivate such a scenario, let us introduce a fictional software developer named Pam.

Pam wants to develop a social media photo-sharing mobile app that analyses her and her friends photos. Pam wants the app to categorise photos into scenes (e.g., day vs. night, landscape vs. indoors), generate brief descriptions of each photo, and catalogue photos of her friends as well as common objects (e.g., all photos with her Border Collie dog, all photos taken on a beach on a sunny day).

Rather than building a computer vision engine from scratch, which would take far too much time and effort, Pam thinks she can achieve this using one of the common computer vision CISs. Pam comes from a typical software engineering background and has insufficient

knowledge of key computer vision terminology and no understanding of the processes behind deep-learning. She ultimately believes all are APIs alike and internalises a deterministic mindset of them; when she decides on one of the three APIs, she expects a static result always. As she expects the same for whenever she calls, for example, any substring API with the call (or similar) of `substring("doggy", 0, 2)` and would expect the response `'dog'` as its output.

To make an assessment of these APIs, she tries her best to read through the documentation of some computer vision APIs, but she has no guiding framework to help her choose the right one. Some of the questions that may come to mind include:

- What does confidence mean? Aren't these APIs consistent?
- Will she need a combination of many computer vision APIs to solve this task?
- How does she know when there is a defect in the response? How can she report it?
- How does she know what labels the API can pick up, and what labels it can't?
- How does she know when the models update? What is the release cycle?
- How does it describe her photos and detect the faces?
- How can she interpret the results if she disagrees with it to help improve her app?

Dazzled by this, she does some brief reading on Wikipedia but is confused by the immense technical detail to take in. She would like some form of guiding framework to assist her and in software engineering terms she can understand.

1.2 Research Goals

In this thesis, we explore the effect stochastic and probabilistic systems play on the usability of APIs with respect to computer vision CISs. Our perspective is software quality—specifically, validation and verification—within such systems and what best practices within the field of software engineering can be applied to assist in operationalisation such systems.

The goals of this study aim to provide a snapshot of current developer best practices towards the usage of CISs to provide a guiding framework and recommendations for software developers and CISs providers alike. We propose two major bodies of work.

Goal 1: ***Understand the developer's mindset towards selecting a computer vision CIS.***

Goal 2: ***Determine what quality factors affect software built using computer vision CISs.***

Goal 3: ***Provide an evaluation framework developers wishing to use computer vision CISs.***

Chiefly, we can specify the following high-level research questions:

RQ1. How do software engineers understand and evaluate computer vision CISs for use in both generic and specific applications?

RQ2. Do software engineers follow best practices when evaluating computer vision CISs? How does this compare to actual practice?

RQ3. What is needed to improve the state-of-the-art of computer vision CISs in terms of API documentation?

RQ4. What aspects of validation and verification can be improved in the field of computer vision CISs?

Ultimately, we seek to understand the conceptual understanding of software engineers who operationalise stochastic and probabilistic systems, and furthermore understand knowledge representation with these systems' API documentation. Our motivation is to provide insight into current practices and compare the best practices with actual practise. We strive for this to provide developers with a guiding framework on how to best operationalise these systems via the form of some checklist or tool they can use to ensure optimal software quality.

It is anticipated that the findings from this study in the computer vision CISs space will be generalisable to other areas, such as time-series information, natural language processing and others.

1.3 Methodology

For this study, we propose running several experiments involving developers and several computer vision CISs, using action-based mixed method approaches and involving documentary analysis. This study will organically evolve by observing phenomena surrounding computer vision API internal quality, chiefly their documentation and responses. We adopt a mixed methods approach, performing both qualitative and quantitative data collection on these two

key aspects by using documentary research methods for inspecting the API documentation and structured observations to quantitatively analyse the results over time (RQs 3 and 4).

Our first proposal for usability studies will survey a number of developers from various levels of seniority and experience (gathering such demographical data to assess a wider sample size) to provide insight into how these developers perceive the non-deterministic nature of computer vision APIs, asking them specific questions about their conceptual understanding of computer vision to identify any outstanding gaps in their knowledge and factor this into known literature (RQs 1 and 2).

We will then conduct a structured interview with a 'mock' computer vision API to remove any developer bias toward any one particular computer vision API that already exists and by which the developer may have already used in the past. Here, we will investigate if developers have any patterns of practice and if they conform to software engineering best practices (RQs 1, 2 and 3).

From these insights, we can then develop a series of assistive recommendations that aide in improving the validation and verification of the existing computer vision API tooling. This may involve a third party tool that helps developers evaluate which particular API is right for their specific computer vision use case.

Chapter 2

Literature Review

⟨ TODO: **Reiterate research claims from Chapter 1 - Introduction.** ⟩

⟨ TODO: **Review literature around this claim from theoretical lenses.** ⟩

2.1 Software Quality

⟨ TODO: **Background on the development of software quality models.** *McCall's model was one of the first software quality models introduced. It described quality from X perspectives... this was further developed by the ISO quality model, which enhanced by Y... In the late 1990s, Dromey's interperation expanded...*⟩

⟨ TODO: **Relate software quality to CV systems; internal & external quality.** ⟩

⟨ TODO: **Discuss gaps in the software quality literature relating directly to CV quality.** ⟩

2.2 Probabilistic and Stochastic Systems

⟨ TODO: **What are stochastic/probabilistic systems? E.g., model interpretation?** ⟩

⟨ TODO: **What understanding might be missing from model interpretation? Relate back to topic.** ⟩

2.2.1 Model Interpretability

As the rise of applied AI increases, the need for engineering interpretability around models becomes paramount. Model interpretability has been stressed since early machine learning research in the late 1980s and 1990s (such as Quinlan [85] and Michie [77]), and although there has since been a significant body of work in the area [98, 19, 87, 26, 93, 71, 23, 60, 18, 47, 32, 105, 21, 40, 70, 74, 81, 106], it is evident that 'accuracy' or model 'confidence' is still

used as a primary criterion for AI evaluation [53, 56, 99]. Indeed, much research into NN or SVM development stresses that 'good' models are those with high accuracy. However, is accuracy enough to justify a model's quality?

To answer this, we revisit what it means for a model to be accurate. Accuracy is an indicator for estimating how well a model's algorithm will work with future or unforeseen data. It is quantified in the AI testing stage, whereby the algorithm is tested against cases known by humans to have ground truth but such cases are unknown by the algorithm. In production, however, all cases are unknown by both the algorithm *and* the humans behind it, and therefore a single value of quality is "not reliable if the future dataset has a probability distribution significantly different from past data" [43], a problem commonly referred to as the *datashift* problem [101]. Analogously, Freitas [43] provides the following description of the problem:

> *The military trained [a NN] to classify images of tanks into enemy and friendly tanks. However, when the [NN] was deployed in the field (corresponding to "future data"), it had a very poor accuracy rate. Later, users noted that all photos of friendly (enemy) tanks were taken on a sunny (overcast) day. I.e., the [NN] learned to discriminate between the colors of the sky in sunny vs. overcast days!* ***If the [NN] had output a comprehensible model (explaining that it was discriminating between colors at the top of the images), such a trivial mistake would immediately be noted.*** *[43]*

So, why must we interpret models? While the formal definition of what it means to be *interpretable* is still somewhat disparate (though some suggestions have been proposed [71]), what is known is (i) there exists a critical trade-off between accuracy and interpretability [42, 59, 62, 49, 36, 111], and (ii) a single quantifiable value cannot satisfy the subjective needs of end-users [43]. As ever-growing domains ML become widespread[1], these applications engage end-users for real-world goals, unlike the aims in early ML research where the aim was to get AI working in the first place. In safety-critical systems where AI provide informativeness to humans to make the final call (see [27, 63, 54]), there is often a mismatch between the formal objectives of the model (e.g., to minimise error) and complex real-world goals, where many other considerations (such as the human factors and cognitive science behind explanations[2])

[1] In areas such as medicine [20, 68, 82, 89, 112, 105, 60, 39, 110, 57, 26], bioinformatics [44, 103, 61, 35, 58], finance [19, 54, 33] and customer analytics [106, 70].

[2] *Interpretations* and *explanations* are often used interchangeably.

are not realised: model optimisation is only worthwhile if they "actually solve the original [human-centred] task of providing explanation" [79] to end-users. **Therefore, when human-decision makers must be interpretable themselves [90], any AI they depend on must also be interpretable.**

Recently, discussion behind such a notion to provide legal implications of interpretability is topical. Doshi-Velez et al. [38] discuss when explanations are not provided from a legal stance—for instance, those affected by algorithmic-based decisions have a 'right to explanation' [48, 107] under the European Union's GDPR[3]. But, explanations are not the only way to ensure AI accountability: theoretical guarantees (mathematical proofs) or statistical evidence can also serve as guarantees [38], however, in terms of explanations, what form they take and how they are proven correct are still open questions [71].

2.2.2 AI Communication Mismatch

From a SE perspective, explanations and interpretability are, by definition, inherently communication issues: what lacks here is a consistent interface between the AI system and the person using it. The ability to encode 'common sense reasoning' [75] into programs today has been achieved, but *decoding* that information is what still remains problematic. At a high level, Shannon and Weaver's theory of communication [97] applies, just as others have done with similar issues in the SE realm [78**?**] (albeit to the domain of visual notations). Humans map the world in higher-level concepts easily when compared to AI systems: while we think of a tree first (not the photons of light or atoms that make up the tree), an algorithm simply sees pixels, and not the concrete object [38] and thusly the AI interprets the tree inversely to humans. Therefore, the interpretation or explanation is done inversely: humans do not explain the individual neurons fired to explain their predictions, and therefore the algorithmic transparent explanations of AI algorithms (*"which neurons were fired to make this AI think this tree is a tree?"*) do not work here.

Therefore, to the user (as mapped using Shannon and Weaver's theory), an AI pipeline (the communication *channel*) begins with a real-world concept, y, that acts as an *information source*. This information source is fed in as a *message*, x, (as pixels) to an AI system (the *transmitter*). The transmitter encodes the pixels to a prediction, $\hat{y}$, the *signal* of the message. This signal is decoded by the *receiver*, an explanation system, $e_x(x, \hat{y})$, that tailors the predic-

[3] `https://www.eugdpr.org` last accessed 13 August 2018.

tion with the given input data to the intended end user (the *destination*) as an explanation, $\tilde{y}$, another type of *message*. Therefore, the user only sees the channel as an input/output pipeline of real-world objects, y, and explanations, $\tilde{y}$, tailored to *them*, without needing to see the inner-mechanics of a prediction $\hat{y}$. We present this diagrammatically in Figure 2.1.

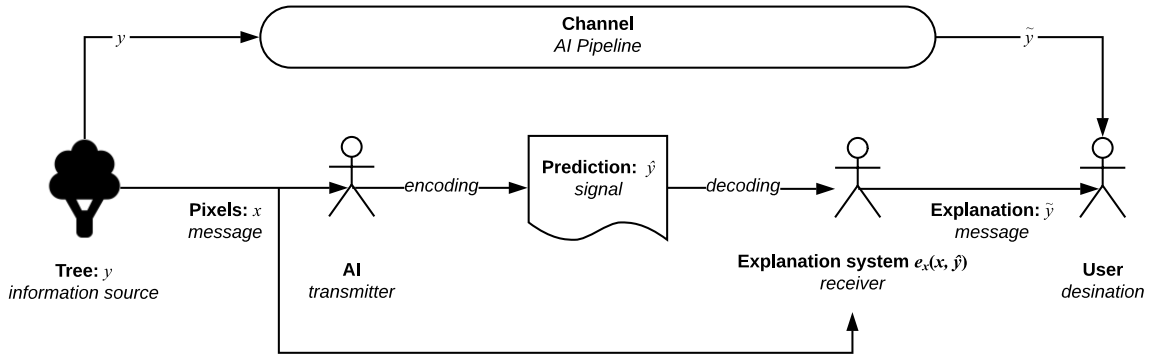

Figure 2.1: Theory of AI communication from information source, y, to intended user as explanations $\tilde{y}$.

2.2.3 Mechanics of Model Interpretation

How do we interpret models? Methods for developing interpretation models include: decision trees [25, 50, 30, 84, 92], decision tables [70**?**] and decision sets [67, 79]; input gradients, gradient vectors or sensitivity analysis [95, 87, 69, 93, 19]; exemplars [64, 45]; generalised additive models [27]; classification (*if-then*) rules [104, 24, 29, 80, 109] and falling rule lists [98]; nearest neighbours [74, 96, 102**?** , 108] and Naïve Bayes analysis [20, 68, 66, 112, 76, 46, 28, 51]. Several cross-domain studies have assessed the interpretability of these techniques against end-users, measuring response time, accuracy in model response and user confidence [54**?** , 17, 100, 94, 44, 74, 106], although it is generally agreed that decision rules and decision tables provide the most interpretation in non-linear models such as SVMs or NNs [44, 74, 106]. For an extensive survey of the benefits and fallbacks of these techniques, we refer to Freitas [43] and Doshi-Velez and Kim [37].

As it stands, AI presents an issue with. (For a detailed discussion, see Doshi-Velez et al. [38].

2.3 Cognitive Biases

⟨ TODO: **Background; what are cognitive biases and how does it relate to SE?** ⟩

⟨ TODO: **Literature of CB specifically in SE.** ⟩

⟨ TODO: **List potential CBs with relation to the AI-based systems.** ⟩

2.4 UX Consistency Principle

⟨ TODO: **Background; what is UX consistency? What does it advocate for and why?** ⟩

⟨ TODO: **What lessons can we learn from UX consistency, and how can we apply it to SE?** ⟩

⟨ TODO: **What are the gaps in SE that do not conform to practices of UX consistency w.r.t. AI systems development?** ⟩

2.5 API Documentation and Standards

⟨ TODO: **What are API documentation standards? What do they advocate for?** ⟩

⟨ TODO: **What is missing for AI documentation? What is the gap?** ⟩

2.6 Validation and Verification

⟨ TODO: **Unsure...** ⟩

2.7 Requirements Specification

⟨ TODO: **Unsure...** ⟩

2.8 Meta-modelling

⟨ TODO: **What is meta-modelling? Can get this from honours thesis...?** ⟩

⟨ TODO: **How does this differ in AI context?** ⟩

Chapter 3

Methodology

3.1 Data Collection and Ethics

3.2 Approach

3.3 Evaluation Methods

3.4 Threats to Validity

3.4.1 Internal Threats

3.4.2 External Threats

3.4.3 Construct

Chapter 4

Project Status

4.1 Completed Work

4.2 Impact

4.3 Timeline

Chapter 5

Conclusion

References

[1] “Amazon Rekognition,” https://aws.amazon.com/rekognition, accessed: 13 September 2018.

[2] “Home - affectiva : Affectiva,” https://www.affectiva.com, accessed: 15 October 2018.

[3] “Image Processing with the Computer Vision API — Microsoft Azure,” https://azure.microsoft.com/en-au/services/cognitive-services/computer-vision/, accessed: 13 September 2018.

[4] “Clarifai,” https://www.clarifai.com, accessed: 13 September 2018.

[5] “Image Recognition API & Visual Search Results,” https://docs.aws.amazon.com/rekognition/latest/dg/labels-detect-labels-image.html, accessed: 13 September 2018.

[6] “The face recognition company - cognitec,” http://www.cognitec.com, accessed: 15 October 2018.

[7] “Image recognition api — deepai,” https://deepai.org/ai-image-processing, accessed: 26 September 2018.

[8] “Vision API - Image Content Analysis — Cloud Vision API — Google Cloud,” https://cloud.google.com/vision/, accessed: 13 September 2018.

[9] “Watson visual recognition,” https://www.ibm.com/watson/services/visual-recognition/, accessed: 13 September 2018.

[10] “Imagga - powerful image recognition apis for automated categorization & tagging in the cloud and on-premises,” https://imagga.com, accessed: 13 September 2018.

[11] “Kairos: Serving businesses with face recognition,” https://www.kairos.com, accessed: 15 October 2018.

[12] “Amazon Mechanical Turk,” https://www.mturk.com, accessed: 15 October 2018.

[13] “Detecting labels in an image,” https://docs.aws.amazon.com/rekognition/latest/dg/labels-detect-labels-image.html, accessed: 13 September 2018.

[14] “Scale: API for Training Data,” https://www.scaleapi.com, accessed: 15 October 2018.

[15] “Image recognition - talkwalker,” https://www.talkwalker.com/image-recognition, accessed: 13 September 2018.

[16] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, and X. Zheng, “TensorFlow: Large-scale machine learning on heterogeneous systems,” 2015, software available from tensorflow.org. [Online]. Available: https://www.tensorflow.org/

[17] H. Allahyari and N. Lavesson, “User-oriented assessment of classification model understandability,” in *11th scandinavian conference on Artificial intelligence*. IOS Press, 2011.

[18] M. G. Augasta and T. Kathirvalavakumar, “Reverse engineering the neural networks for rule extraction in classification problems,” *Neural processing letters*, vol. 35, no. 2, pp. 131–150, 2012.

[19] D. Baehrens, T. Schroeter, S. Harmeling, M. Kawanabe, K. Hansen, and K.-R. MÃžller, “How to explain individual classification decisions,” *Journal of Machine Learning Research*, vol. 11, no. Jun, pp. 1803–1831, 2010.

[20] R. Bellazzi and B. Zupan, “Predictive data mining in clinical medicine: current issues and guidelines,” *International journal of medical informatics*, vol. 77, no. 2, pp. 81–97, 2008.

[21] A. Ben-David, “Monotonicity maintenance in information-theoretic machine learning algorithms,” *Machine Learning*, vol. 19, no. 1, pp. 29–43, 1995.

[22] J. J. Blake, L. P. Maguire, B. Roche, T. M. McGinnity, and L. J. McDaid, "The Implementation of Fuzzy Systems, Neural Networks and Fuzzy Neural Networks using FPGAs." *Inf. Sci.*, 1998.

[23] O. Boz, "Extracting decision trees from trained neural networks," in *Proceedings of the eighth ACM SIGKDD international conference on Knowledge discovery and data mining*. ACM, 2002, pp. 456–461.

[24] M. Bramer, *Principles of data mining*. Springer, 2007, vol. 180.

[25] L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone, *Classification and Regression Trees*. CRC press, 1984.

[26] A. Bussone, S. Stumpf, and D. O'Sullivan, "The Role of Explanations on Trust and Reliance in Clinical Decision Support Systems." *ICHI*, 2015.

[27] R. Caruana, Y. Lou, J. Gehrke, P. Koch, M. Sturm, and N. Elhadad, "Intelligible Models for HealthCare - Predicting Pneumonia Risk and Hospital 30-day Readmission." *KDD*, pp. 1721–1730, 2015.

[28] J. Cheng and R. Greiner, "Learning bayesian belief network classifiers: Algorithms and system," in *Conference of the Canadian Society for Computational Studies of Intelligence*. Springer, 2001, pp. 141–151.

[29] P. Clark and R. Boswell, "Rule induction with CN2: Some recent improvements," in *European Working Session on Learning*. Springer, 1991, pp. 151–163.

[30] M. Craven and J. W. Shavlik, "Extracting Tree-Structured Representations of Trained Networks." *NIPS*, 1995.

[31] H. da Mota Silveira and L. C. Martini, "How the New Approaches on Cloud Computer Vision can Contribute to Growth of Assistive Technologies to Visually Impaired in the Following Years," *Journal of Information Systems Engineering and Management*, vol. 2, no. 2, pp. 1–3, 2017.

[32] K. Dejaeger, F. Goethals, A. Giangreco, L. Mola, and B. Baesens, "Gaining insight into student satisfaction using comprehensible data mining techniques," *European Journal of Operational Research*, vol. 218, no. 2, pp. 548–562, 2012.

[33] V. Dhar, D. Chou, and F. Provost, "Discovering Interesting Patterns for Investment Decision Making with GLOWER—A Genetic Learner Overlaid with Entropy Reduction," *Data Mining and Knowledge Discovery*, vol. 4, no. 4, pp. 251–280, 2000.

[34] V. C. Dibia, M. Ashoori, A. Cox, and J. D. Weisz, "TJBot," in *the 2017 CHI Conference Extended Abstracts*. New York, New York, USA: ACM Press, 2017, pp. 381–384.

[35] M. Doderer, K. Yoon, J. Salinas, and S. Kwek, "Protein subcellular localization prediction using a hybrid of similarity search and error-correcting output code techniques that produces interpretable results," *In silico biology*, vol. 6, no. 5, pp. 419–433, 2006.

[36] P. Domingos, "Occam's two razors: The sharp and the blunt," in *KDD*, 1998, pp. 37–43.

[37] F. Doshi-Velez and B. Kim, "Towards a rigorous science of interpretable machine learning," 2017.

[38] F. Doshi-Velez, M. Kortz, R. Budish, C. Bavitz, S. Gershman, D. O'Brien, S. Schieber, J. Waldo, D. Weinberger, and A. Wood, "Accountability of AI Under the Law: The Role of Explanation," *arXiv.org*, Nov. 2017.

[39] W. Elazmeh, W. Matwin, D. O'Sullivan, W. Michalowski, and W. Farion, "Insights from predicting pediatric asthma exacerbations from retrospective clinical data," in *Evaluation Methods for Machine Learning II–Papers from 2007 AAAI Workshop*, 2007, pp. 10–15.

[40] A. J. Feelders, "Prior knowledge in economic applications of data mining," in *European Conference on Principles of Data Mining and Knowledge Discovery*. Springer, 2000, pp. 395–400.

[41] R. T. Fielding, "*Architectural Styles and the Design of Network-based Software Architectures* ," Ph.D. dissertation, University of California, Irvine.

[42] A. A. Freitas, "A critical review of multi-objective optimization in data mining: a position paper," *ACM SIGKDD Explorations Newsletter*, vol. 6, no. 2, pp. 77–86, 2004.

[43] ——, "Comprehensible classification models," *ACM SIGKDD Explorations Newsletter*, vol. 15, no. 1, pp. 1–10, Mar. 2014.

[44] A. A. Freitas, D. C. Wieser, and R. Apweiler, "On the importance of comprehensible classification models for protein function prediction," *IEEE/ACM Transactions on Computational Biology and Bioinformatics (TCBB)*, vol. 7, no. 1, pp. 172–182, 2010.

[45] B. J. Frey and D. Dueck, "Clustering by Passing Messages Between Data Points," *Science*, vol. 315, no. 5814, pp. 972–976, Feb. 2007.

[46] N. Friedman, D. Geiger, and M. Goldszmidt, "Bayesian network classifiers," *Machine Learning*, vol. 29, no. 2-3, pp. 131–163, 1997.

[47] G. Fung, S. Sandilya, and R. B. Rao, "Rule extraction from linear support vector machines," in *Proceedings of the eleventh ACM SIGKDD international conference on Knowledge discovery in data mining*. ACM, 2005, pp. 32–40.

[48] B. Goodman and S. R. Flaxman, "EU regulations on algorithmic decision-making and a "right to explanation"." *CoRR*, 2016.

[49] P. D. Grünwald, *The minimum description length principle*. MIT press, 2007.

[50] T. Hastie, R. Tibshirani, and J. H. Friedman, *The Elements of Statistical Learning*, ser. Data Mining, Inference, and Prediction. Springer Science & Business Media, Jan. 2001.

[51] D. Heckerman, D. M. Chickering, C. Meek, R. Rounthwaite, and C. Kadie, "Dependency networks for inference, collaborative filtering, and data visualization," *Journal of Machine Learning Research*, vol. 1, no. Oct, pp. 49–75, 2000.

[52] C. Howard. (2018, May) Introducing Google AI. [Online]. Available: https://ai.googleblog.com/2018/05/introducing-google-ai.html

[53] J. Huang and C. X. Ling, "Using AUC and accuracy in evaluating learning algorithms," *IEEE Transactions on knowledge and Data Engineering*, vol. 17, no. 3, pp. 299–310, 2005.

[54] J. Huysmans, K. Dejaeger, C. Mues, J. Vanthienen, and B. Baesens, "An empirical evaluation of the comprehensibility of decision table, tree and rule based predictive models," *Decision Support Systems*, vol. 51, no. 1, pp. 141–154, Apr. 2011.

[55] A. Iyengar, "Supporting Data Analytics Applications Which Utilize Cognitive Services," in *2017 IEEE 37th International Conference on Distributed Computing Systems (ICDCS)*. IEEE, May 2017, pp. 1856–1864.

[56] N. Japkowicz and M. Shah, *Evaluating learning algorithms: a classification perspective*. Cambridge University Press, 2011.

[57] M. W. M. Jaspers, M. Smeulers, H. Vermeulen, and L. W. Peute, "Effects of clinical decision-support systems on practitioner performance and patient outcomes: a synthesis of high-quality systematic review findings," *Journal of the American Medical Informatics Association*, vol. 18, no. 3, pp. 327–334, 2011.

[58] T. Jiang and A. E. Keating, "AVID: an integrative framework for discovering functional relationships among proteins," *BMC bioinformatics*, vol. 6, no. 1, p. 136, 2005.

[59] Y. Jin, *Multi-objective machine learning*. Springer Science & Business Media, 2006, vol. 16.

[60] U. Johansson and L. Niklasson, "Evolving decision trees using oracle guides," in *IEEE Symposium on Computational Intelligence and Data Mining (CIDM)*. IEEE, 2009.

[61] A. Karwath and R. D. King, "Homology induction: the use of machine learning to improve sequence similarity searches," *BMC bioinformatics*, vol. 3, no. 1, p. 11, 2002.

[62] K. A. Kaufman and R. S. Michalski, "Learning from inconsistent and noisy data: the AQ18 approach," in *International Symposium on Methodologies for Intelligent Systems*. Springer, 1999, pp. 411–419.

[63] B. Kim, *Interactive and Interpretable Machine Learning Models for Human Machine Collaboration*. Massachusetts Institute of Technology, 2015.

[64] B. Kim, C. Rudin, and J. A. Shah, "The Bayesian Case Model - A Generative Approach for Case-Based Reasoning and Prototype Classification." *NIPS*, 2014.

[65] A. J. Ko and Y. Riche, "The role of conceptual knowledge in API usability," in *2011 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*. IEEE, 2011, pp. 173–176.

[66] I. Kononenko, "Inductive and Bayesian learning in medical diagnosis," *Applied Artificial Intelligence an International Journal*, vol. 7, no. 4, pp. 317–337, 1993.

[67] H. Lakkaraju, S. H. Bach, and J. Leskovec, "Interpretable Decision Sets - A Joint Framework for Description and Prediction." *KDD*, pp. 1675–1684, 2016.

[68] N. Lavrač, "Selected techniques for data mining in medicine," *Artificial intelligence in medicine*, vol. 16, no. 1, pp. 3–23, 1999.

[69] T. Lei, R. Barzilay, and T. Jaakkola, "Rationalizing Neural Predictions," *arXiv.org*, Jun. 2016.

[70] E. Lima, C. Mues, and B. Baesens, "Domain knowledge integration in data mining using decision tables: Case studies in churn prediction," *Journal of the Operational Research Society*, vol. 60, no. 8, pp. 1096–1106, 2009.

[71] Z. C. Lipton, "The Mythos of Model Interpretability." *CoRR*, 2016.

[72] D. Lo Giudice, C. Mines, A. LeClair, R. Curran, and A. Homan, "How AI Will Change Software Development And Applications," Tech. Rep., Nov. 2016.

[73] T. E. Marshall and S. L. Lambert, "Cloud-based intelligent accounting applications: accounting task automation using IBM watson cognitive computing," *Journal of Emerging Technologies in Accounting*, vol. 15, no. 1, pp. 199–215, 2018.

[74] D. Martens, J. Vanthienen, W. Verbeke, and B. Baesens, "Performance of classification models from a user perspective," *Decision Support Systems*, vol. 51, no. 4, pp. 782–793, 2011.

[75] J. McCarthy, "Programs with Common Sense," Cambridge, MA, USA, Tech. Rep., 1960.

[76] D. Michie, D. J. Spiegelhalter, and C. C. Taylor, "Machine Learning and Statistical Classification of Artificial intelligence," 1994.

[77] D. Michie, "Machine Learning in the Next Five Years." *EWSL*, 1988.

[78] D. L. Moody, "The "Physics" of Notations - Toward a Scientific Basis for Constructing Visual Notations in Software Engineering." *IEEE Trans. Software Eng.*, 2009.

[79] M. Narayanan, E. Chen, J. He, B. Kim, S. Gershman, and F. Doshi-Velez, "How do Humans Understand Explanations from Machine Learning Systems? An Evaluation of the Human-Interpretability of Explanation." *CoRR*, 2018.

[80] F. E. Otero and A. A. Freitas, "Improving the interpretability of classification rules discovered by an ant colony algorithm," in *Proceedings of the 15th annual conference on Genetic and evolutionary computation*. ACM, 2013, pp. 73–80.

[81] M. Pazzani, "Comprehensible knowledge discovery: gaining insight from data," in *First Federal Data Mining Conference and Exposition*, 1997, pp. 73–82.

[82] M. J. Pazzani, S. Mani, and W. R. Shankle, "Acceptance of rules generated by machine learning among medical experts," *Methods of information in medicine*, vol. 40, no. 05, pp. 380–385, 2001.

[83] M. Piccioni, C. A. Furia, and B. Meyer, "An Empirical Study of API Usability," in *2013 ACM / IEEE International Symposium on Empirical Software Engineering and Measurement*. IEEE, 2013, pp. 5–14.

[84] J. R. Quinlan, "C4. 5: Programming for machine learning," *Morgan Kauffmann*, vol. 38, p. 48, 1993.

[85] ——, "Some elements of machine learning," in *International Conference on Inductive Logic Programming*. Springer, 1999, pp. 15–18.

[86] A. Reis, D. Paulino, V. Filipe, and J. Barroso, "Using Online Artificial Vision Services to Assist the Blind - an Assessment of Microsoft Cognitive Services and Google Cloud Vision." *WorldCIST*, vol. 746, no. 12, pp. 174–184, 2018.

[87] M. T. Ribeiro, S. Singh, and C. Guestrin, ""Why Should I Trust You?"," in *the 22nd ACM SIGKDD International Conference*. New York, New York, USA: ACM Press, 2016, pp. 1135–1144.

[88] M. Ribeiro, K. Grolinger, and M. A. M. Capretz, "MLaaS: Machine Learning as a Service," in *2015 IEEE 14th International Conference on Machine Learning and Applications (ICMLA)*. IEEE, Dec. 2015, pp. 896–902.

[89] G. Richards, V. J. Rayward-Smith, P. H. Sönksen, S. Carey, and C. Weng, "Data mining for indicators of early mortality in a database of clinical records," *Artificial intelligence in medicine*, vol. 22, no. 3, pp. 215–231, 2001.

[90] G. Ridgeway, D. Madigan, T. Richardson, and J. O'Kane, "Interpretable Boosted Naïve Bayes Classification." *KDD*, 1998.

[91] M. P. Robillard, "What makes APIs hard to learn? Answers from developers," *IEEE Software*, vol. 26, no. 6, pp. 27–34, 2009.

[92] L. Rokach and O. Z. Maimon, *Data mining with decision trees: theory and applications*. World scientific, 2008, vol. 69.

[93] A. S. Ross, M. C. Hughes, and F. Doshi-Velez, "Right for the Right Reasons: Training Differentiable Models by Constraining their Explanations," *arXiv.org*, Mar. 2017.

[94] M. Schwabacher, P. Langley, and P. Norvig, "Discovering communicable scientific knowledge from spatio-temporal data," *ICML*, 2001.

[95] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, "Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization," in *2017 IEEE International Conference on Computer Vision (ICCV)*. IEEE, 2017, pp. 618–626.

[96] S. Sen and L. Knight, "A genetic prototype learner," in *IJCAI*. Citeseer, 1995, pp. 725–733.

[97] C. E. Shannon and W. Weaver, *The mathematical theory of communication*. Urbana, IL: The University of Illinois Press, 1963.

[98] S. Singh, M. T. Ribeiro, and C. Guestrin, "Programs as Black-Box Explanations," *arXiv.org*, Nov. 2016.

[99] M. Sokolova and G. Lapalme, "A systematic analysis of performance measures for classification tasks," *Information Processing & Management*, vol. 45, no. 4, pp. 427–437, 2009.

[100] G. H. Subramanian, J. Nosek, S. P. Raghunathan, and S. S. Kanitkar, "A comparison of the decision table and tree," *Communications of the ACM*, vol. 35, no. 1, pp. 89–94, 1992.

[101] M. Sugiyama, N. D. Lawrence, and A. Schwaighofer, *Dataset shift in machine learning*. The MIT Press, 2017.

[102] N. R. Suri, V. S. Srinivas, and M. N. Murty, "A cooperative game theoretic approach to prototype selection," in *European Conference on Principles of Data Mining and Knowledge Discovery*. Springer, 2007, pp. 556–564.

[103] D. Szafron, P. Lu, R. Greiner, D. S. Wishart, B. Poulin, R. Eisner, Z. Lu, J. Anvik, C. Macdonell, and A. Fyshe, "Proteome Analyst: custom predictions with explanations in a web-based tool for high-throughput proteome annotations," *Nucleic acids research*, vol. 32, no. 2, pp. W365–W371, 2004.

[104] S. Thrun, "Is Learning The n-th Thing Any Easier Than Learning The First?" p. 7, 1996.

[105] A. Van Assche and H. Blockeel, "Seeing the forest through the trees: Learning a comprehensible model from an ensemble," in *European conference on machine learning*. Springer, 2007, pp. 418–429.

[106] W. Verbeke, D. Martens, C. Mues, and B. Baesens, "Building comprehensible customer churn prediction models with advanced rule induction techniques," *Expert Systems with Applications*, vol. 38, no. 3, pp. 2354–2364, 2011.

[107] S. Wachter, B. Mittelstadt, and L. Floridi, "Why a Right to Explanation of Automated Decision-Making Does Not Exist in the General Data Protection Regulation," *International Data Privacy Law*, vol. 7, no. 2, pp. 76–99, Jun. 2017.

[108] D. Wettschereck, D. W. Aha, and T. Mohri, "A review and empirical evaluation of feature weighting methods for a class of lazy learning algorithms," *Artificial Intelligence Review*, vol. 11, no. 1-5, pp. 273–314, 1997.

[109] I. H. Witten, E. Frank, M. A. Hall, and C. J. Pal, *Data Mining: Practical machine learning tools and techniques*. Morgan Kaufmann, 2016.

[110] M. L. Wong and K. S. Leung, *Data mining using grammar based genetic programming and applications*. Springer Science & Business Media, 2006, vol. 3.

[111] J. Zahálka and F. Železný, "An experimental test of Occam's razor in classification," *Machine Learning*, vol. 82, no. 3, pp. 475–481, 2011.

[112] B. Zupan, J. DemšAr, M. W. Kattan, J. R. Beck, and I. Bratko, “Machine learning for survival analysis: a case study on recurrence of prostate cancer,” *Artificial intelligence in medicine*, vol. 20, no. 1, pp. 59–75, 2000.